



```
In [2]: # Import necessary Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Load the dataset
dataset_path = "C:\\Users\\ARYAN PARIKH\\Desktop\\Oasis Internship\\archive (2
car_data = pd.read_csv(dataset_path)

# Display the first few rows of the dataset to understand its structure
print(car_data.head())

# Explore the dataset to understand the features and target variable
print(car_data.info())

# Select relevant features and target variable
X = car_data[['Year', 'Present_Price', 'Driven_kms', 'Fuel_Type', 'Selling_type']]
y = car_data['Selling_Price']

# Convert categorical variables into dummy/indicator variables
X = pd.get_dummies(X, drop_first=True)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a Random Forest Regressor model
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared Score: {r2}')

# Visualize predicted vs. actual prices
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Selling Price')
plt.ylabel('Predicted Selling Price')
plt.title('Actual vs. Predicted Selling Price')
plt.show()
```

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	\
0	ritz	2014	3.35	5.59	27000	Petrol	
1	sx4	2013	4.75	9.54	43000	Diesel	
2	ciaz	2017	7.25	9.85	6900	Petrol	
3	wagon r	2011	2.85	4.15	5200	Petrol	
4	swift	2014	4.60	6.87	42450	Diesel	

	Selling_type	Transmission	Owner
0	Dealer	Manual	0
1	Dealer	Manual	0
2	Dealer	Manual	0
3	Dealer	Manual	0
4	Dealer	Manual	0

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 301 entries, 0 to 300
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Car_Name	301 non-null	object
1	Year	301 non-null	int64
2	Selling_Price	301 non-null	float64
3	Present_Price	301 non-null	float64
4	Driven_kms	301 non-null	int64
5	Fuel_Type	301 non-null	object
6	Selling_type	301 non-null	object
7	Transmission	301 non-null	object
8	Owner	301 non-null	int64

```
dtypes: float64(2), int64(3), object(4)
```

```
memory usage: 21.3+ KB
```

```
None
```

```
Mean Squared Error: 0.8209857844262285
```

```
R-squared Score: 0.9643601062650229
```



In [ ]: